

Locking to ensure serializability

- ✓ *Concurrent access* to database items is controlled by strategies based on *locking*, *timestamping* or *certification*
- ✓ A *lock* is an *access privilege to a single database item*
- ✓ *Lock Manager*: manages the locks requested by transactions.
- ✓ Locks are
 - ✓ *obtained by transactions*
 - ✓ stored in a *lock table*
 - ✓ Lock is an entry of the form *(item, lock-type, transactionID)*
 - ✓ *item* is the item that the transaction locks
 - ✓ *lock-type* can be *shared* or *exclusive*
 - ✓ *transactionID* is the transaction identifier

Locking to Ensure Serializability

- ✓ Locking can *prevent the lost update problem*
 - ✓ $T_1 = \text{Lock}(A) R_1(A) W_1(A) \text{Unlock}(A) C_1$
 - ✓ $T_2 = \text{Lock}(A) R_2(A) W_2(A) \text{Unlock}(A) C_2$
- ✓ Under the *locking policy*, *only serial execution of the transactions is permitted*

Locking

- ✓ When a transaction holds an *exclusive* lock on a database item, *no other transaction can read or write the item*
 - ✓ used for writing
- ✓ When a transaction holds a *shared* lock, *other transactions can obtain a shared lock on the same item*
 - ✓ used for reading
- ✓ **Assumptions**
 - ✓ there is a *single type of lock* and
 - ✓ *every transaction must obtain a lock* on an item before accessing it.
 - ✓ *all items locked by a transaction must be unlocked, otherwise no other transaction may gain access to them.*
 - ✓ a transaction *must wait* until the lock it requests is released by the transaction that holds it.

Transaction Management

1. *Locking* can prevent the *lost update* problem:

$$T_1 = \text{Lock}_1(A) \ R_1(A) \ W_1(A) \ \text{Unlock}_1(A) \ C_1$$

$$T_2 = \text{Lock}_2(A) \ R_2(A) \ W_2(A) \ \text{Unlock}_2(A) \ C_2$$

2. *Locking enforces a serial execution of the transactions*

3. Locking can also prevent the *blind write* problem:

$$T_1 = \text{Lock}_1(A) \ W_1(A) \ \text{Lock}_1(B) \ W_1(B) \ \text{Unlock}_1(A) \ \text{Unlock}_1(B) \ C_1$$

$$T_2 = \text{Lock}_2(A) \ W_2(A) \ \text{Lock}_2(B) \ W_2(B) \ \text{Unlock}_2(A) \ \text{Unlock}_2(B) \ C_1$$

✓ Then the following schedule is valid:

Lock₁(A) *W₁(A)* *Lock₁(B)* *W₁(B)* *Unlock₁(A)* *Lock₂(A)* *W₂(A)*
Unlock₁(B) *Lock₂(B)* *W₂(B)* *Unlock₂(A)* *Unlock₂(B)* *C₁ C₂*

LiveLock

- ✓ Undesirable phenomena if *locks are granted in an arbitrary manner*
- ✓ **Example:**
 - ✓ while T2 is waiting for T1 to release the lock on A, another transaction T3 that has also requested a lock on A is granted the lock instead of T2. When T3 releases the lock on A the lock is granted to T4 etc.
- ✓ *Livelock: The situation where a transaction may wait for ever while other transactions obtain a lock on a database item*
 - ✓ Can be avoided by using a first-come-first-served lock granting strategy but, even then a *deadlock* might occur

Deadlock

- ✓ Occurs *when a transaction is waiting to lock an item that is currently locked by some other transaction*
- ✓ **Example:** Consider the transactions:
$$T_1 = \text{Lock}_1(A) \text{ Lock}_1(B) \dots \text{Unlock}_1(A) \text{ Unlock}_1(B) C_1$$
$$T_2 = \text{Lock}_2(B) \text{ Lock}_2(A) \dots \text{Unlock}_2(B) \text{ Unlock}_1(A) C_2$$
- ✓ Assume T_1 is *granted* a *lock on A* and T_2 is *granted* a *lock on B*
- ✓ Then T_1 *requests* a *lock on B* but is **forced to wait** because T_2 *has* the *lock on B*.
- ✓ Similarly, T_2 *requests* a *lock on A* but is **forced to wait** because T_1 has the *lock on A*.


Neither transaction can proceed because each is waiting for the other to release a lock: both processes wait for ever

Different solutions for Deadlocks

- ✓ **Solution 1:** Require *each transaction to request all locks at once. Either all locks are granted or none.*
- ✓ **Solution 2:** Assign *an arbitrary linear order* to the *items* and *require all transactions to request their locks in that order.*
- ✓ **Solution 3:** *Do nothing to prevent deadlocks: abort* one or more of the deadlocked transactions if a deadlock arises.

Deadlock Discovery

- ✓ Deadlocks can be discovered using *wait-for graphs*:
 - ✓ Given a set of transactions S , a *wait-for graph* is a *directed graph*:
 - ✓ vertices correspond to transactions in the set
 - ✓ there exists an edge from T_i to T_j if T_i is waiting to lock an item on which T_j is holding a lock.
- ✓ **Theorem:** A set of transactions is deadlocked if and only if there exists a cycle in the *wait-for graph*.
- ✓ **Example:** The *wait-for graph* for the transactions contains a cycle
$$T_1 = \text{Lock}_1(A) \text{ Lock}_1(B) \dots \text{Unlock}_1(A) \text{ Unlock}_1(B) C_1$$
$$T_2 = \text{Lock}_2(B) \text{ Lock}_2(A) \dots \text{Unlock}_2(B) \text{ Unlock}_2(A) C_2$$



2-Phase Locking (2PL)

- ✓ 2-Phase Locking (2PL): a protocol ensuring *serializability of schedules*
- ✓ *Definition:* A schedule is said to obey the *2-phase locking* protocol if the following rules are obeyed by *each transaction* in the schedule
 1. When a transaction attempts to *read* (*write*) a data item, a *read lock* (*write lock*) must be acquired first
 2. If a transaction T_1 holds a lock on data item A for operation op_1 and some other transaction T_2 requests the lock to perform a *conflicting operation* op_2 on the same item, the transaction requesting the lock (T_2) is forced to *wait until no conflicting lock on the item exists*
 - ✓ *(only read locks are non-conflicting)*
 - 3 *A transaction cannot request additional locks once it releases any lock!*

2-Phase Locking (2PL): Conflicts

- ✓ two locks by the *same transaction never conflict*
- ✓ a transaction with a *read lock on a data item* can acquire a *write lock on the item as long as no other transaction has a lock on the data item;*
- ✓ a transaction with *a write lock on a data item need not acquire a read lock* on the same item.
- ✓ *2PL permits the early release of locks*
- ✓ Notation:
 - ✓ RL_i : transaction T_i obtains a *read* lock
 - ✓ WL_i : transaction T_i obtains a *write* lock
 - ✓ RU_i : transaction T_i releases a *read* lock
 - ✓ WU_i : transaction T_i releases a *write* lock

2-Phase Locking (2PL): Example

- ✓ Does the following schedule obey the 2PL protocol?

$S = R_1(A) R_2(B) W_2(B) R_2(A) W_2(A) R_1(B) C_1 C_2$

Lock/unlock operations must be *added first*. The schedule becomes:

$S' = \boxed{RL_1(A)} R_1(A) RU_1(A) \boxed{RL_2(B)} R_2(B) \boxed{WL_2(B)} W_2(B) WU_2(B) \boxed{RL_2(A)}$
 $R_2(A) \boxed{WL_2(A)} W_2(A) \boxed{RL_1(B)} R_1(B) C_1 C_2$

- ✓ *Rule 1* : no item is accessed without a lock being granted to the requested transaction
- ✓ *obeyed*

2-Phase Locking (2PL): Example

- ✓ Does the following schedule obey the 2PL protocol?

$S = R_1(A) R_2(B) W_2(B) R_2(A) W_2(A) R_1(B) C_1 C_2$

Lock/unlock operations must be *added first*. The schedule becomes:

$S = \boxed{RL_1(A)} \text{ } \underline{R_1(A)} \text{ } \boxed{RU_1(A)} \text{ } RL_2(B) \text{ } R_2(B) \text{ } WL_2(B) \text{ } W_2(B) \text{ } WU_2(B) \text{ } RL_2(A)$
 $\text{ } R_2(A) \text{ } \boxed{WL_2(A)} \text{ } \underline{W_2(A)} \text{ } RL_1(B) \text{ } R_1(B) \text{ } C_1 \text{ } C_2$

- ✓ *Rule 2* : no two conflicting operations have a lock on the same item at the same time
- ✓ *obeyed*

2-Phase Locking (2PL): Example

- ✓ Does the following schedule obey the 2PL protocol?

$S = R_1(A) R_2(B) W_2(B) R_2(A) W_2(A) R_1(B) C_1 C_2$

- ✓ *Lock/unlock operations* must be *added first*. The schedule becomes:

$S = RL_1(A) R_1(A) RU_1(A) RL_2(B) R_2(B) WL_2(B) W_2(B) WU_2(B) RL_2(A) R_2(A)$
 $WL_2(A) W_2(A) RL_1(B) R_1(B) C_1 C_2$

- ✓ *Rule 3: A transaction cannot request additional locks once it releases any lock!*
 - ✓ *Violated!*

2-Phase Locking (2PL): Example

- ✓ Applying the 2PL discipline to the schedule

$S = R_1(A) R_2(B) W_2(B) R_2(A) W_2(A) R_1(B) C_1 C_2$

yields the following interleaved execution (all locks released at commit)

T_1	$RL_1(A)$	$R_1(A)$							
T_2			$RL_2(B)$	$R_2(B)$	$WL_2(B)$	$W_2(B)$	$RL_2(A)$	$R_2(A)$	$WL_2(A)$



T_1		$RL_1(B)$	<i>wait</i>	<i>abort</i>			<i>restart</i>	C_1
T_2	<i>wait</i>				$W_2(A)$	C_2		

The deadlock had to be resolved by aborting and restarting one of the transactions.

Under 2PL S is *equivalent* to the *serial schedule* $T_2 T_1$

2-Phase Locking (2PL): Example

- ✓ **Theorem:** A schedule that follows 2PL *is always serializable*.
- ✓ **Example:**
 - ✓ The schedule $S' = R_1(A) R_2(A) W_1(A) W_2(A) C_1 C_2$ is forced to execute as follows by a transaction scheduler that uses 2PL:

$T1$	$RL_1(A)$	$R_1(A)$			$WL_1(A)$	<i>wait</i>			<i>abort</i>
$T2$			$RL_2(A)$	$R_2(A)$			$WL_2(A)$	<i>wait</i>	

$T1$			<i>restart</i>	C_1
$T2$	$W_2(A)$	C_2		

If no locking was imposed S' would be non-serializable

Transaction Management

- ✓ **Example:** Consider the following transactions
 - ✓ $T_1: W_1(U) R_1(Y) W_1(U) C_1$
 - ✓ $T_2: R_2(X) W_2(U) W_2(Y) W_2(W) C_2$
 - ✓ $T_3: W_3(W) R_3(X) W_3(U) W_3(Z) C_3$
- ✓ Is it possible to add lock/unlock steps to these transactions so that every legal schedule is serializable?
- ✓ Answer: *yes by adding add lock/unlock steps using 2PL*

Transaction Management

1. $T_1: W_1(U) R_1(Y) W_1(U) C_1$
2. $T_2: R_2(X) W_2(U) W_2(Y) W_2(W) C_2$
3. $T_3: W_3(W) R_3(X) W_3(U) W_3(Z) C_3$

T_1							
T_2					$RL_2(X)$	$WL_2(U)$	$wait$
T_3	$WL_3(W)$	$RL_3(X)$	$WL_3(U)$	$WL_3(Z)$			$W_3(W)$

T_1							$WL_1(U)$
T_2	$wait$	$wait$	$wait$	$wait$	$wait$	$WL_2(Y)$	
T_3	$WU_3(W)$	$R_3(X)$	$RU_3(X)$	$W_3(U)$	$WU_3(U)$		

Transaction Management

- 1. $T_1: W_1(U) R_1(Y) W_1(U) C_1$
- 2. $T_2: R_2(X) W_2(U) W_2(Y) W_2(W) C_2$
- 3. $T_3: W_3(W) R_3(X) W_3(U) W_3(Z) C_3$

